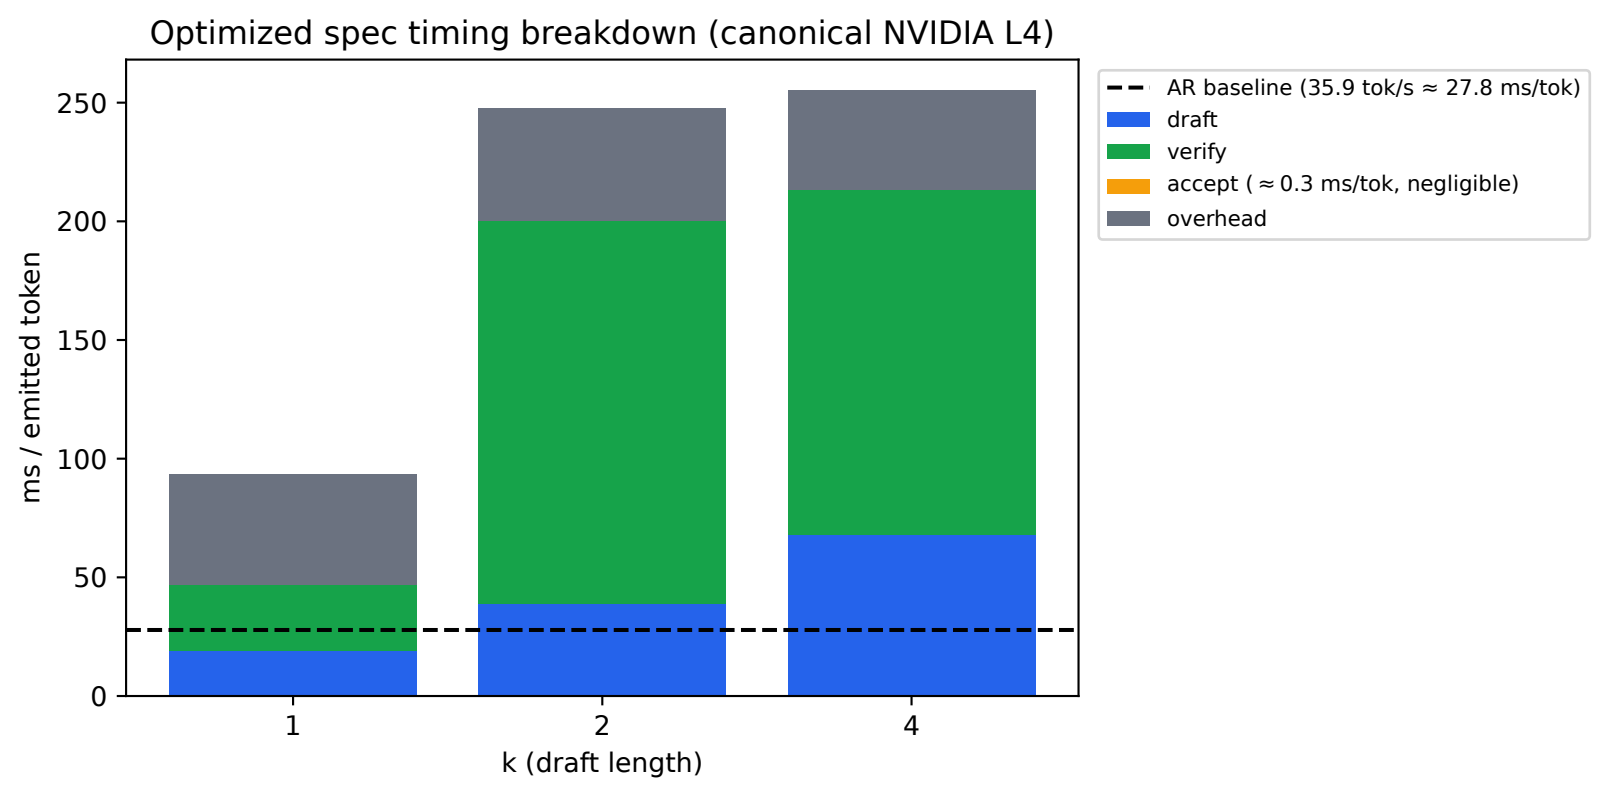
Optimized spec timing breakdown (canonical NVIDIA L4)
ms / emitted token
0
50
100
150
200
250
1
2
4
k (draft length)
AR baseline (35.9 tok/s ≈ 27.8 ms/tok)
draft
verify
accept (≈ 0.3 ms/tok, negligible)
overhead